



Monsoon 2021

Positional Index

- **Creating Inverted Index with Term Positions**

Dr. Rajendra Prasath

Indian Institute of Information Technology Sri City, Chittoor



27th August 2021 (rajendra.2power3.com)

> Topics to be covered

- Recap:
 - Inverted Index Construction
 - Term - Document Matrix

- Boolean Operators
- Boolean Retrieval
- Boolean Queries

- Text Collection / Corpora
- Evaluation Strategy

- More topics to come up ... Stay tuned ...!!

Recap: Information Retrieval

- **Information Retrieval (IR)** is finding material (usually documents) of an unstructured nature (usually text) that satisfies an information need from within large collections (usually stored on computers).
- These days we frequently think first of web search, but there are many other cases:
 - E-mail search
 - Searching your laptop
 - Corporate knowledge bases
 - Legal information retrieval
 - and so on . . .

Recap: Boolean Incidence Matrix


Terms	d_1	d_2	d_3	...	d_n
the	1	1	1	...	0
a	1	1	1	...	1
Darjeeling	1	1	1	...	0
is	1	1	1	...	0
of	1	1	1	...	0
in	1	0	0	...	1
and	1	1	0	...	0
Bengal	1	0	1	...	0
It	1	0	1	...	0
Its	0	1	0	...	1
state	1	0	1	...	0
West	1	0	1	...	1

Boolean Queries: Exact match

- ✧ The Boolean retrieval model is being able to ask a query that is a Boolean expression:
 - ✧ Boolean Queries are queries using AND, OR and NOT to join query terms
 - ✧ Views each document as a set of words
 - ✧ Is precise: document matches condition or not
 - ✧ Perhaps the simplest model to build an IR system on
- ✧ Primary commercial retrieval tool for 3 decades
- ✧ Many search systems you still use are Boolean:
 - ✧ Email, library catalog, Mac OS X Spotlight

Query optimization

- ✧ What is the best order for query processing?
- ✧ Consider a query that is an AND of n terms.
- ✧ For each of the n terms, get its postings, then AND them together.



Query: **Brutus AND Calpurnia AND Caesar**

Query Processing - Exercises

- ✧ Exercise: If the query is *friends AND romans AND (NOT countrymen)*, how could we use the freq of *countrymen*?
- ✧ Exercise: Extend the merge to an arbitrary Boolean query. Can we always guarantee execution in time linear in the total postings size?
- ✧ Hint: Begin with the case of a Boolean *formula* query: in this, each query term appears only once in the query

Phrase queries

- We want to be able to answer queries such as “stanford university” – as a phrase
- Thus the sentence “**I went to university at Stanford**” is not a match.
 - The concept of phrase queries has proven easily understood by users; one of the few “advanced search” ideas that works
 - Many more queries are *implicit phrase queries*
- For this, it no longer suffices to store only
 - $\langle term : docs \rangle$ entries

A First attempt: Biword indexes

- Index **every consecutive pair of terms** in the text as a phrase
- For example the text “Friends, Romans, Countrymen” would generate the **biwords**
 - **friends romans**
 - **romans countrymen**
- Each of these **biwords** is now a dictionary term
- Two-word phrase query-processing is now immediate

Longer phrase queries

- Longer phrases can be processed by breaking them down
- ***stanford university palo alto*** can be broken into the Boolean query on biwords:

stanford university AND university palo AND palo alto

Without the docs, we cannot verify that the docs matching the above Boolean query do contain the phrase

Can have false positives!

Issues for Biword Indexes

- False positives, as noted before
- Index blowup due to bigger dictionary
 - Infeasible for more than Biwords, big even for them
- Biword indexes are not the standard solution (for all Biwords) but can be part of a compound strategy

Solution 2: Positional indexes

- Positional indexes are a more efficient alternative to biword indexes
- Postings lists in a nonpositional index:
 - each posting is just a docID
- Postings lists in a positional index:
 - each posting is a docID and a list of positions

- In the postings, store, for each **term** the position(s) in which tokens of it appear:

<**term**, number of docs containing **term**;

doc1: position1, position2 ... ;

doc2: position1, position2 ... ;

etc.>

Solution 2: Positional indexes

Positional index example

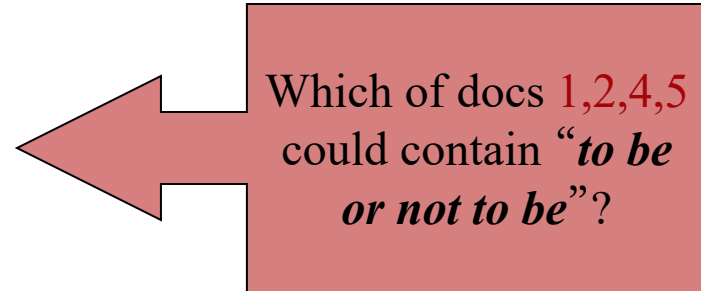
<*be*: 993427;

1: 7, 18, 33, 72, 86, 231;

2: 3, 149;

4: 17, 191, 291, 430, 434;

5: 363, 367, ...>



- For phrase queries, we use a merge algorithm recursively at the document level
- But we now need to deal with more than just equality

Processing a phrase query

- Extract inverted index entries for each distinct term: ***to, be, or, not***
- Merge their *doc:position* lists to enumerate all positions with “***to be or not to be***”
 - ***to:***
 - 2:1,17,74,222,551; 4:8,16,190,429,433; 7:13,23,191; ...
 - ***be:***
 - 1:17,19; 4:17,191,291,430,434; 5:14,19,101; ...
- Same general method for proximity searches

Positional Indexes: Example

Query: “ $to_1 be_2 or_3 not_4 to_5 be_6$ ”

TO, 993427:

1: $\langle 7, 18, 33, 72, 86, 231 \rangle$;

2: $\langle 1, 17, 74, 222, 255 \rangle$;

4: $\langle 8, 16, 190, 429, 433 \rangle$;

5: $\langle 363, 367 \rangle$;

7: $\langle 13, 23, 191 \rangle$; ...

BE, 178239:

1: $\langle 17, 25 \rangle$;

4: $\langle 17, 191, 291, 430, 434 \rangle$;

5: $\langle 14, 19, 101 \rangle$; ...

Document 4 is a match!

Positional index size

- A positional index expands postings storage substantially
- Even though indices can be compressed
- Nevertheless, a positional index is now standardly used because of the power and usefulness of phrase and proximity queries
- used explicitly or implicitly in a ranking retrieval system

Positional index size

- Need an entry for each occurrence, not just once per document
- Index size depends on average document size
 - Average web page has < 1000 terms
 - Books, even some epic poems ... easily 100,000 terms
- Consider a term with frequency 0.1%

Why?

Document size	Postings	Positional postings
1000	1	1
100,000	1	100

Rules of thumb

- ✧ A positional index is 2–4 as large as a non-positional index
- ✧ Positional index size 35–50% of volume of original text
 - Caveat: all of this holds for “English-like” languages

Proximity queries

- ✧ LIMIT! /3 STATUTE /3 FEDERAL /2 TORT
 - ✧ Again, here, / k means “**within k words of**”
- ✧ Clearly, positional indexes can be used for such queries; biword indexes cannot.
- ✧ **Exercise:** Adapt the linear merge of postings to handle proximity queries
- ✧ Can you make it work for any value of k ?
 - ✧ This is a little tricky to do correctly and efficiently

Proximity Search

- ✧ We can also use it for proximity search.
For example: employment /4 place
- ✧ Find all documents that contain EMPLOYMENT and PLACE within 4 words of each other
- ✧ Employment agencies that place healthcare workers are seeing growth is a hit
- ✧ Employment agencies that have learned to adapt now place healthcare workers is not a hit

Proximity Search

- ✧ Use Positional Index
- ✧ Simplest Algorithm: look at cross-product of positions of
 - ✧ EMPLOYMENT in document and
 - ✧ PLACE in document
- ✧ Very inefficient for frequent words, especially stop words
- ✧ Note that we want to return the actual matching positions, not just a list of documents
- ✧ This is important for dynamic summaries, etc

Combination schemes

- These two approaches can be profitably combined
 - For particular phrases (“**Michael Jackson**”, “**Britney Spears**”) it is inefficient to keep on merging positional postings lists
 - Even more so for phrases like “**The Who**”
- Williams et al. (2004) evaluate a more sophisticated mixed indexing scheme
 - A typical web query mixture was executed in $\frac{1}{4}$ of the time of using just a positional index
 - It required 26% more space than having a positional index alone

Combination Scheme (contd.)

- ✧ Biword indexes and positional indexes can be profitably combined
- ✧ Many biwords are extremely frequent: Michael Jackson, Britney Spears etc
- ✧ For these biwords, increased speed compared to positional postings intersection is substantial.
- ✧ Combination scheme:
 - ✧ Include frequent biwords as vocabulary terms in the index
 - ✧ Do all other phrases by positional intersection

Summary

In this class, we focused on:

- (a) Understanding of the basic unit of classical information retrieval systems
- (b) Positional Indexes
 - i. Positional Index Size
- (c) Proximity Search
- (d) Combination Schemes



Questions It's Your Time

